

Review Session 10

Sophie Hill

04/08/2022

PSet 3

Last one!

Q1-Q10: Difference-in-differences (RS9)

Protip: Read the intro and skim the Card & Krueger paper!

Q11-Q14:

- Interpreting coefficients (RS7)
- Omitted variable bias (RS8)
- Fixed effects (RS9)

Q15-Q17: Predicting a numeric outcome (RS10)

Q18-Q20: Predicting a binary outcome, logistic regression (RS10)

Agenda



Logistic regression

- Why?
- Interpreting coefficients
- Predicted probabilities

Prediction

- Classification threshold
- Accuracy, precision, recall
- For numeric outcomes: RMSE, MAE, MAPE

Logistic regression: Why?



Logistic regression: R code

```
# Logistic regression
glm(vote_dem ~ age + white,
    data = cces20,
    family = "binomial") %>%
tidy()
```

```
## # A tibble: 3 × 5
##   term          estimate std.error statistic  p.value
##   <chr>         <dbl>     <dbl>     <dbl>    <dbl>
## 1 (Intercept)    2.07      0.0394      52.5  0
## 2 age          -0.0189    0.000629   -30.1 2.10e-198
## 3 white        -0.858     0.0250   -34.3 9.43e-258
```

How is this different to linear regression with `lm()`?

- `glm()` instead of `lm()`
- `family = "binomial"` tells R that the outcome variable is binary
- interpretation of the coefficients is different

Interpreting logistic coefficients

Recall that the logit function is defined by:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

So this is our regression equation:

$$\text{logit}\left[\widehat{P(\text{vote_dem} = 1)}\right] = 2.069 - 0.019(\text{age}) - 0.858(\text{white})$$


Or, equivalently:

$$\widehat{P(\text{vote_dem} = 1)} = \text{logit}^{-1}\left[2.069 - 0.019(\text{age}) - 0.858(\text{white})\right]$$

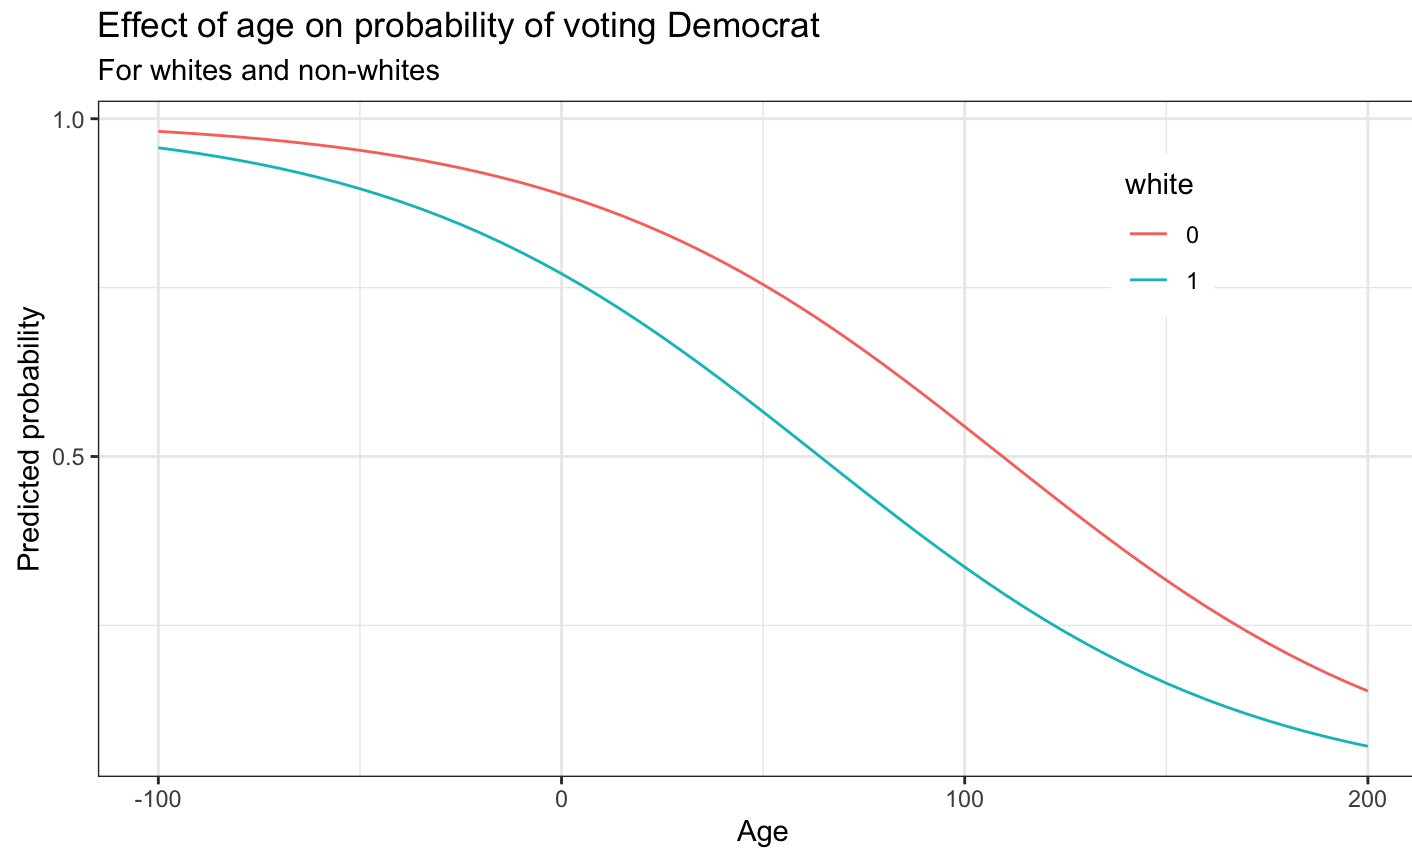
The S-curve

Effect of age on probability of voting Democrat

For whites and non-whites



The S-curve



The “divide-by-4” trick

$$\text{logit}\left[\widehat{P(\text{vote_dem} = 1)}\right] = 2.069 - 0.019(\text{age}) - 0.858(\text{white})$$

Useful trick: Divide a logit coefficient by 4 to get an upper bound of the ppt change in $P(Y_i = 1)$.

$$-0.019/4 = -0.005 \qquad -0.858/4 = -0.215$$

- Increasing **age** by 1 year is associated with no more than a 0.5ppt decrease in the probability of voting democratic.
- Being **white** (vs non-white) is associated with no more than a 22ppt decrease in the probability of voting democratic.

Caution: this is just a rough rule-of-thumb. It is more accurate when the outcome probability is close to 0.5 (where the S-curve is the steepest).

Change in predicted probabilities

$$\text{logit}\left[\widehat{P(\text{vote_dem} = 1)}\right] = 2.069 - 0.019(\text{age}) - 0.858(\text{white})$$

The divide-by-4 rule gives us a rough upper bound. To interpret the coefficient on X_i more precisely, we should compute predicted probabilities at different values of X_i .

The effect of a 1-unit change in X_i can vary depending on the values of the other predictors. (This is a non-linear model!) So we also need to pick some sensible values to hold the other predictors constant at.

Change in predicted probabilities

```
ggpredict(mod_logistic,  
          terms = c("age [20:80, by = 10]",  
                    "white [1]"))
```

```
## # Predicted probabilities of vote_dem
```

```
##
```

```
## age | Predicted |          95% CI
```

```
## -----
```

```
## 20 |          0.70 | [0.69, 0.71]
```

```
## 30 |          0.66 | [0.65, 0.66]
```

```
## 40 |          0.61 | [0.61, 0.62]
```

```
## 50 |          0.57 | [0.56, 0.57]
```

```
## 60 |          0.52 | [0.51, 0.52]
```

```
## 70 |          0.47 | [0.46, 0.48]
```

```
## 80 |          0.43 | [0.42, 0.43]
```

Change in predicted probabilities

```
ggpredict(mod_logistic,  
          terms = c("white"))
```

```
## # Predicted probabilities of vote_dem
```

```
##
```

```
## white | Predicted |          95% CI
```

```
## -----
```

```
##      0 |          0.74 | [0.74, 0.75]
```

```
##      1 |          0.55 | [0.55, 0.56]
```

```
##
```

```
## Adjusted for:
```

```
## * age = 53.14
```

Classification

```
augment(mod_logistic,  
        cces20,  
        type.predict = "response") %>%  
  select(1:4) %>%  
  head()
```

```
## # A tibble: 6 × 4  
##   age white vote_dem .fitted  
##   <dbl> <dbl>   <dbl>   <dbl>  
## 1    54     1     0    0.547  
## 2    74     1     1    0.453  
## 3    58     1     1    0.529  
## 4    59     1     0    0.524  
## 5    73     1     0    0.458  
## 6    50     1     0    0.566
```

Choosing the threshold

```
cces20 <-  
  augment(mod_logistic, cces20,  
          type.predict = "response") %>%  
  select(1:4) %>%  
  mutate(pred = ifelse(.fitted > 0.5, 1, 0))
```

```
head(cces20)
```

```
## # A tibble: 6 × 5
```

```
##   age white vote_dem .fitted pred  
##   <dbl> <dbl>   <dbl>   <dbl> <dbl>  
## 1    54     1       0  0.547     1  
## 2    74     1       1  0.453     0  
## 3    58     1       1  0.529     1  
## 4    59     1       0  0.524     1  
## 5    73     1       0  0.458     0  
## 6    50     1       0  0.566     1
```

How good are our predictions?

Accuracy: out of all predictions, what % are correct?

$$\frac{TP+TN}{TP+FP+TN+FN}$$

Precision: out of all *predicted* 1's, what % are correct?

$$\frac{TP}{TP+FP}$$

Recall: out of all *actual* 1's, what % are correctly predicted?


$$\frac{TP}{TP+FN}$$

How good are our predictions?

```
cces20 %>%  
  summarize(  
    accuracy = sum(vote_dem == pred) / nrow(cces20),  
    precision = sum(vote_dem == 1 & pred == 1) / sum(pred == 1),  
    recall = sum(vote_dem == 1 & pred == 1) / sum(vote_dem == 1)  
  )
```

```
## # A tibble: 1 × 3  
##   accuracy precision recall  
##   <dbl>      <dbl>  <dbl>  
## 1    0.607    0.631  0.823
```


How good are our predictions?




Intuition check

For low thresholds, recall is 1. For high thresholds, recall is 0. **Why?**

```
cces20 %>% select(.fitted) %>% summary()
```

```
##      .fitted  
##  Min.      :0.3577  
## 1st Qu.:0.5050  
##  Median :0.5753  
##   Mean   :0.5967  
## 3rd Qu.:0.6808  
##   Max.   :0.8492
```

Intuition check



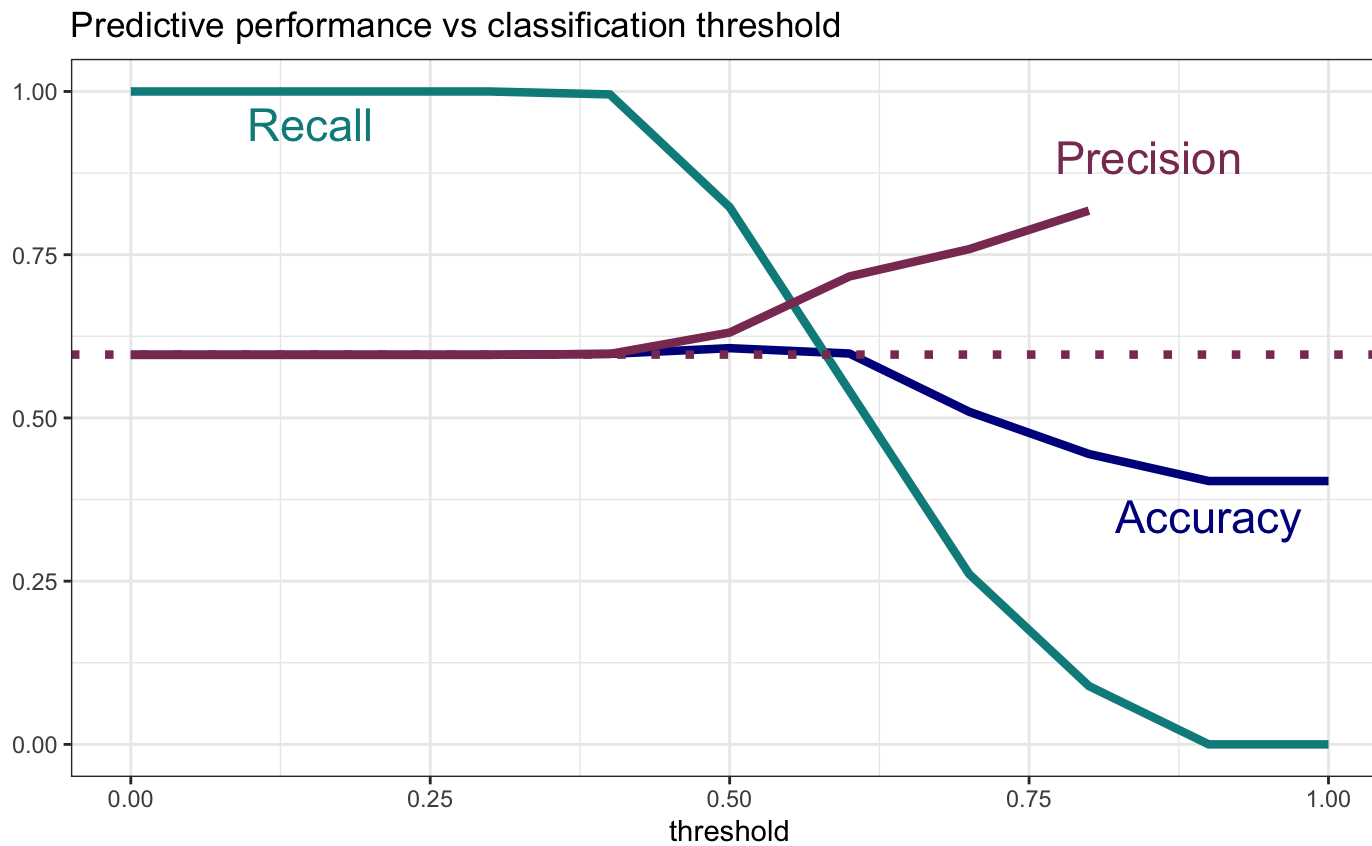
Intuition check

For low thresholds, accuracy and precision are constant at 0.597. **Why?**

```
cces20 %>%  
  summarize(mean_y = mean(vote_dem))
```

```
## # A tibble: 1 × 1  
##   mean_y  
##   <dbl>  
## 1  0.597
```

Intuition check



Predicting a numeric variable

Let's try to predict someone's *age* based on their marital status.

```
mod_age1 <- cces %>%  
  lm(age ~ marstat,  
     data = .)
```

```
tidy(mod_age1)
```

```
## # A tibble: 3 × 5  
##   term                estimate std.error statistic  p.value  
##   <chr>              <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)        57.5      0.205     281.    0  
## 2 marstatMarried    -4.66     0.227    -20.6 9.36e-94  
## 3 marstatSingle    -20.4     0.245    -83.1 0
```

Predicting a numeric variable

```
cces <- augment(mod_age1, cces) %>%  
  select(age:.fitted) %>%  
  mutate(error = .fitted - age)
```

```
head(cces)
```

```
## # A tibble: 6 × 7
```

```
##      age faminc marstat  ownhome pid3  .fitted  error  
##    <dbl> <dbl> <chr>    <fct>   <chr>   <dbl>   <dbl>  
## 1     54  15000 Married  Rent    Rep     52.9   -1.13  
## 2     65  65000 Divorced Rent    Dem     57.5   -7.47  
## 3     58 110000 Married  Own     Dem     52.9   -5.13  
## 4     53  15000 Single   Rent    Ind     37.1  -15.9  
## 5     59  45000 Divorced Own     Rep     57.5   -1.47  
## 6     50  55000 Married  Own     Rep     52.9    2.87
```

Measures of predictive accuracy

Root Mean Squared Error (RMSE)

$$\sqrt{\frac{1}{N} \sum_i (Y_i - \hat{Y}_i)^2}$$

Mean Absolute Error (MAE)

$$\frac{1}{N} \sum_i |Y_i - \hat{Y}_i|$$

Mean Absolute Percent Error (MAPE)

$$\frac{1}{N} \sum_i 100 \cdot \frac{|Y_i - \hat{Y}_i|}{Y_i}$$

Measures of predictive accuracy

We can calculate the RMSE, MAE, and MAPE manually:

```
cces %>%
  summarize(rmse = sqrt(sum(error^2) / nrow(cces)),
            mae = sum(abs(error)) / nrow(cces),
            mape = sum((abs(error)/age) * 100) / nrow(cces)
  )

## # A tibble: 1 × 3
##   rmse   mae  mape
##   <dbl> <dbl> <dbl>
## 1  15.0  12.5  30.7
```

Measures of predictive accuracy

Or we can use the canned functions `rmse()`, `mae()`, `mape()` in the `{tidymodels}` package:

```
cces %>%
  summarize(rmse = rmse(cces, truth = age, estimate = .fitted)$estimate,
            mae = mae(cces, truth = age, estimate = .fitted)$estimate,
            mape = mape(cces, truth = age, estimate = .fitted)$estimate
  )

## # A tibble: 1 × 3
##   rmse   mae  mape
##   <dbl> <dbl> <dbl>
## 1  15.0  12.5  30.7
```

Improving our model

What if we include home ownership (`ownhome`) and household income (`faminc`) in our model too?

```
mod_age2 <- cces %>%  
  lm(age ~ marstat + ownhome + faminc,  
      data = .)
```

```
tidy(mod_age2)
```

```
## # A tibble: 5 × 5
```

##	term	estimate	std.error	statistic	p.value
##	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	(Intercept)	54.8	0.226	242.	0
## 2	marstatMarried	-5.74	0.233	-24.6	4.90e-133
## 3	marstatSingle	-19.5	0.240	-81.0	0
## 4	ownhomeOwn	7.55	0.169	44.7	0
## 5	faminc	-0.0000274	0.00000177	-15.5	2.90e- 54

Do you think the RMSE, MAE, MAPE will be smaller or larger?

Comparing two models

Initial model with just marital status as a predictor:

```
## # A tibble: 1 × 3
##   rmse   mae  mape
##   <dbl> <dbl> <dbl>
## 1  15.0  12.5  30.7
```

New model with home ownership and family income added:

```
## # A tibble: 1 × 3
##   rmse   mae  mape
##   <dbl> <dbl> <dbl>
## 1  14.6  12.2  29.7
```

Test vs training set

Of course, if we were really interested in the predictive power of our model we should look at how it performs on *new data* - i.e. data it was not trained on.

```
cces <-  
  cces %>%  
  mutate(test = sample(  
    x = c(0, 1),  
    size = nrow(cces),  
    replace = TRUE  
  ))  
  
cces_split <- cces %>% group_by(test) %>% group_split()  
cces_train <- cces_split[[1]]  
cces_test <- cces_split[[2]]
```

Test vs training set

```
mod_age2 <- cces_train %>%  
  lm(age ~ marstat + ownhome + faminc,  
      data = .)  
  
cces_test %>%  
  mutate(.fitted = predict(mod_age2, .),  
         error = .fitted - age) %>%  
  summarize(  
    rmse = rmse(., truth = age, estimate = .fitted)$estimate,  
    mae = mae(., truth = age, estimate = .fitted)$estimate,  
    mape = mape(., truth = age, estimate = .fitted)$estimate  
  )  
  
## # A tibble: 1 × 3  
##   rmse   mae  mape  
##   <dbl> <dbl> <dbl>  
## 1  14.6  12.2  29.7
```